

Kalman Folding 9: in C (WORKING DRAFT)

Extracting Models from Data, One Observation at a Time

Brian Beckman

<2016-05-17 Tue>

Contents

1	Abstract	1
2	Literate Code	1
3	This Document is Literate	2
3.1	Get C to Work	2
3.2	Recurrence for the Mean	2
3.3	FoldList and Recurrence for the Variance	6
4	Basic Kalman Folding	10
4.1	Get GSL to Work	10
4.2	BLAS Through GSL	12
4.3	Step 2: Getting LAPACK to work	13

1 Abstract

In this literate program, we show one way to implement the basic Kalman fold in C. The background for Kalman folding appears in *Kalman Folding, Part 1*,¹ where we highlight the the unique advantages of a functional fold for deploying test-hardened code verbatim in harsh, mission-critical environments. In that and in other papers in this series, the emphasis is on mathematical clarity and conciseness. Here, the emphasis is on embedded systems by using the primary language for implementing such systems, C.

2 Literate Code

A *literate program*² is a single text document that contains both

- \LaTeX typesetting instructions and

¹B. Beckman, *Kalman Folding, Part 1*, to appear.

²https://en.wikipedia.org/wiki/Literate_programming

- source code and scripts for an application.

A literate program can be edited in any text editor and controlled with text-based versioning tools like git, but requires other tools to extract the typeset document and to extract the source code and lay it out on disk as required by compilers, linkers, and other build tools. My chosen tool is emacs and org-mode.³ It allows interactive evaluation of code and you may wish to follow along interactively. To do that, you will need emacs and the text source for this document. The text source for this document is found here [TODO](#). A distribution of emacs for the Mac with adequate org-babel support is maintained by Vincent Goulet at the University of Laval.⁴ Emacs for Linux and Windows is easy to find, but I do not discuss it further. *Spacemacs*⁵ offers vim emulation with the tools needed for literate programming.

3 This Document is Literate

3.1 Get C to Work

First, make sure the following works⁶ in org-babel and org-babel tangle. If it does work, you have a correctly installed C compiler.

```
int a=7;
int b=7;
printf("%d\n", a*b);
// ~~> produces
```

49

3.2 Recurrence for the Mean

3.2.1 Business Code

C code is usually much longer than the corresponding Wolfram code (or Lisp, or Haskell, or OCaml, etc.). The extra length, overall, is due mostly to the need for manual memory management in C. However, we also don't have pattern matching for unpacking inputs and we don't have literal list / array notation for expressing some inputs and for packing outputs.

The goal of our C code is to make the final result as clean and as close to the original design in Wolfram as possible. For instance, the original source for *Recurrence for the Mean* in paper 1¹ is

```
cume[{x_, n_}, z_] := (* pattern matching for unpacking inputs *)
  With[{K = 1/(n + 1)},
    {x + K (z - x), n + 1}]; (* literal notation for packing outputs *)
Fold[cume, {0, 0}, {55, 89, 144}] (* literal notation for some inputs *)
~~> {96, 3}
```

³<https://en.wikipedia.org/wiki/Org-mode>

⁴<http://vgoulet.act.ulaval.ca/en/emacs/>

⁵<http://www.spacemacs.org>

⁶Make sure the first example from <http://tinyurl.com/kz2lz7m> works

The equivalent “business” code in C is still longer, but can see that the code for computing the gain and for the resulting *Accumulation* is virtually identical (we must use an explicit multiplication operator `*`, whereas a space suffices in Wolfram for scalar multiplication). At least the memory management is hidden in the types *Accumulation* and *Observation*, articulated further below.

```
{ Accumulator cume = ^(Accumulation a, Observation z)
    { /* unpack inputs */
      T x = a.elements[0];
      T n = a.elements[1];
      /* compute gain */
      T K = 1.0 / (1.0 + n);
      /* business logic, and packing results */
      Accumulation r;
      r.elements[0] = x + K * (z - x);
      r.elements[1] = 1.0 + n;

      return r;    };

  Accumulation x0 = zeroAccumulation ();

  Observation tmp[3] = {55, 89, 144};
  Observations zs = createObservations(3, tmp);

  Accumulation result = fold (cume, x0, zs);

  printAccumulation (result);
  return 0;    }
```

Figure 1: Business logic in C

3.2.2 Primary Numerical Type

We abstract out the primary numerical type into the traditional `T` to make it easier to change.

```
typedef double T;
```

Figure 2: Primary numerical type

3.2.3 Accumulation Type

The code for the *Accumulation* type is a bit elaborate, but the extra abstractions will serve us well when we get to the Kalman filter.

The *Accumulation* structure presumes that all values are copied around on every use, and that’s safe, and also means that we don’t need `alloc` & `free` routines for this type. These accumulation types are usually small, so the time needed to copy them around may be acceptable. More

sophisticated memory management for them entails more code, so we opt for keeping the code small at the cost of some copying that could be optimized away.

Also, in the interest of saving space, specifically, staircases of closing curly braces on lines by themselves, we adopt the *Pico*⁷ style for bracing.

```
const size_t Accumulation_size = 3;
typedef struct s_Accumulation
{   T elements[Accumulation_size];   } Accumulation, * pAccumulation;

Accumulation zeroAccumulation (void)
{   Accumulation r;
    memset ((void *)r.elements, 0, Accumulation_size * sizeof (T));
    return r;   }

void printAccumulation (Accumulation a)
{   printf ("{" );
    for (size_t i = 0; i < Accumulation_size; ++i)
    {   printf ("%lf", a.elements[i]);
        if (i < Accumulation_size - 1)
        {   printf (", ");   }   }
    printf ("}\n");   }
```

Figure 3: Accumulation type

We have harmlessly used 3 for the accumulation size because we want to reuse this code later. We could make it variable at the cost of more unilluminating code.

3.2.4 Observation Types

Because we don't statically know the number of observations, we must use dynamic memory allocation. In an embedded application, we would use arena memory (fixed-length circular buffer pools of fixed-length structs) or stack allocation (*calloc*). Here, for brevity and because this is a testing deployment, we use heap memory (stdlib's *malloc* and *free*). These are unacceptable in embedded applications because of fragmentation and unbounded execution times.

When we get to lazy streams, we won't need these at all. They're only for arrays of observations all in memory at one time.

The primary helper type is a bounded array of *Observations* type that includes the length and a handy iterator-like *current* index. Most of the code for this type concerns explicit memory management for this helper type.

We also include an *Observation* type, for asbstraction hygiene.

```
typedef T Observation, * pObservation;
typedef struct s_BoundedArray_Observations
{   int count;
    int current;
    pObservation observations;   } Observations;
```

⁷<http://tinyurl.com/gku2k74>

```

/*private*/pObservation allocObservationArray (int count_)
{
    /* Don't use malloc & free in embedded apps. Use arena or stack memory. */
    pObservation po = (pObservation) malloc (count_ * sizeof (Observation));
    if (NULL == po)
    {
        printf ("Failed to alloc %d observations\n", count_);
        exit (-1);
    }
    return po;
}

Observations createObservations (int count_, pObservation pObservations)
{
    pObservation po = allocObservationArray (count_);
    memcpy ((void *)po, (void *)pObservations, sizeof (Observation) * count_);
    Observations result;
    result.count = count_;
    result.current = 0;
    result.observations = po;
    return result;
}

void freeObservations (Observations o)
{
    /* Don't use malloc & free in embedded apps. Use arena or stack memory. */
    free ((void *)o.observations);
}

```

Figure 4: Observation types

3.2.5 Accumulator Type

Our last type definition is for the *Accumulator* function. Here we cheat a bit and use an extension to the C language called *Blocks*,⁸ which implements full closures. We could explicitly implement enough of closures for our purposes, but this extension is widely available with clang and llvm on Apple computers and Linux, and it's too convenient to pass up. With compilers for bare-metal processors in embedded systems, we might not have it and have to do more work by hand. With this extension, the *Accumulator* type, defined with the hat syntax ^, behaves just like a function pointer, which would be defined with the ordinary pointer syntax, *.

```
typedef Accumulation (^Accumulator) (Accumulation a, Observation b);
```

Figure 5: Accumulator type

3.2.6 The Fold Over Observations

The final piece is the *fold* operator. This particular one knows details of the *Observations* type, so is specific to it. We have another fold over lazy streams, articulated below, just as with Wolfram.

```
Accumulation fold (Accumulator f, Accumulation x0, Observations zs)
```

⁸<http://tinyurl.com/bgwfkyk>

```

{   for (zs.current = 0; zs.current < zs.count; ++zs.current)
    {   x0 = f (x0, zs.observations[zs.current]);   }
    return x0;   }

```

Figure 6: Fold over observations in bounded arrays

3.2.7 Pulling it All Together

```

#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <Block.h>
<<c-numerical-type>>
<<c-accumulation-type>>
<<c-observation-types>>
<<c-accumulator-type>>
<<c-fold-over-observations>>
int main (int argc, char ** argv)
<<c-business-logic>>

```

Figure 7: Recurrence for the mean: entire program

Tangle this code out to a C file by executing ‘org-babel-tangle’ while visiting this literate source code in emacs.

Compile and run the code as follows:

```

gcc -Wall -Werror recurrenceForTheMean.c -o recurrenceForTheMean
./recurrenceForTheMean

```

Figure 8: Build and execute script for recurrence-for-the-mean

Table 1: Output of recurrence-for-the-mean
{96.000000 3.0 0.000000}

producing results all-but-identical to those from the Wolfram language.

3.3 FoldList and Recurrence for the Variance

The original paper introduced Wolfram’s *FoldList* along with the recurrence for the variance. We do likewise here, implementing our own *foldList* in C.

3.3.1 Bounded Array for Accumulations

FoldList produces a list of accumulations, one for the initial accumulation and another for each observation. With lists of observations all in memory, we could calculate the length of the output and

preallocate a list of accumulations of the correct size, but we are not able to do that with lazy streams of observations or asynchronous observables of observations. We opt, then, for on-demand, dynamic memory management for the output accumulations. “On-demand,” here, means growing the output array as new accumulations arrive. We use the common trick of doubling the capacity of the output array every time the capacity is exceeded. This trick is a reasonable compromise of space and time efficiency.

We emulate the *bounded-array* interface created for observations, and add three more functions to the usual *create*, *free*, and *print*.

lastAccumulations returns the last accumulation in a bounded array; needed for *foldList*

appendAccumulations appends a new accumulation to a bounded array of accumulations, growing the capacity if needed

foldList takes an accumulator *f*, an initial accumulation a_0 , a bounded array of observations *zs*, and produces a bounded array of accumulations.

```
typedef struct s_BoundedArray_Accumulations
{
    int count;
    int max;
    pAccumulation accumulations ;    } Accumulations;

Accumulation lastAccumulations (Accumulations as)
{
    if (0 == as.count)
    {
        printf ("Attempt to pull non-existent element\n");
        exit (-4);    }
    return as.accumulations[as.count - 1];    }

Accumulations appendAccumulations (Accumulations as, Accumulation a)
{
    Accumulations result = as;
    if (result.count + 1 > result.max)
    {
        /* Double the storage. */
        int new_max = 2 * result.max;
        /* Don't use malloc & free in embdded apps. Use arena or stack memory. */
        pAccumulation new = (pAccumulation)
            malloc (sizeof (Accumulation) * new_max);
        if (NULL == new)
        {
            printf ("Failed to alloc %d Accumulations\n", new_max);
            exit (-2);    }
        if (result.count != result.max)
        {
            printf ("Internal bugcheck\n");
            exit (-3);    }
        memset ((void *)new, 0, new_max * sizeof (Accumulation));
        memcpy ((void *)new, (void *)result.accumulations,
            (sizeof (Accumulation) * result.max));
        free ((void *) result.accumulations);
        result.accumulations = new;
        result.max = new_max;    }
```

```

    result.accumulations[result.count] = a;
    ++ result.count;
    return result;    }

Accumulations createAccumulations (void)
{
    Accumulations result;
    const int init_size = 4;
    result.max = init_size;
    result.count = 0;
    result.accumulations = (pAccumulation)
        malloc (sizeof (Accumulation) * init_size);
    memset ((void *)result.accumulations, 0,
        sizeof (Accumulation) * init_size);
    return result;    }

void freeAccumulations (Accumulations as)
{
    memset ((void *) as.accumulations, 0,
        (sizeof (Accumulation) * as.count));
    free ((void *) as.accumulations);    }

void printAccumulations (Accumulations as)
{
    for (int j = 0; j < as.count; ++j )
        { printAccumulation (as.accumulations[j]);    }    }

Accumulations foldList (Accumulator f, Accumulation a0, Observations zs)
{
    Accumulations result = createAccumulations ();
    result = appendAccumulations (result, a0);
    for (zs.current = 0; zs.current < zs.count; ++zs.current)
        { result = appendAccumulations (
            result,
            f(lastAccumulations(result),
                zs.observations[zs.current]));    }
    return result;    }

```

Figure 9: Bounded array for accumulations

3.3.2 Pulling Together Recurrence for the Variance

```

#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <Block.h>
<<c-numerical-type>>
<<c-accumulation-type>>
<<c-observation-types>>
<<c-accumulator-type>>

```



```

<<c-fold-over-observations>>
<<c-bounded-array-for-accumulations>>
int main (int argc, char ** argv)
{
  Observation tmp[3] = {55, 89, 144};
  Observations zs = createObservations(3, tmp);
  Accumulation x0 = zeroAccumulation ();
  Accumulator cume = ^(Accumulation a, Observation z)
    {
      T var = a.elements[0];
      T x    = a.elements[1];
      T n    = a.elements[2];

      T K = 1.0 / (1.0 + n);
      T x2 = x + K * (z - x);
      T ssr2 = (n - 1.0) * var + K * n * (z - x) * (z - x);

      Accumulation r;
      r.elements[0] = ssr2 / (n > 1.0 ? n : 1.0);
      r.elements[1] = x2;
      r.elements[2] = n + 1.0;
      return r;    };

  Accumulations results = foldList (cume, x0, zs);
  printAccumulations (results);

  freeAccumulations (results);
  freeObservations (zs);
  return 0;    }

```

Figure 10: Recurrence for the variance: entire program

```

gcc -Wall -Werror recurrenceForTheVariance.c -o recurrenceForTheVariance
./recurrenceForTheVariance

```

Figure 11: Build and execute script for recurrence-for-the-variance

Table 2: Output of recurrence-for-the-variance

{0.000000	0.0	0.000000}
{0.000000	55.0	1.000000}
{578.000000	72.0	2.000000}
{2017.000000	96.0	3.000000}

This result is semantically identical to that produced by the following Wolfram code:

```

cume[{var_, x_, n_}, z_] :=
  With[{K = 1/(n + 1)},

```

```

With[{x2 = x + K (z - x),
      ssr2 = (n - 1) var + K n (z - x)^2},
{ssr2/Max[1, n], x2, n + 1}]];
Fold[cume, {0, 0, 0}, zs]
~~> {2017, 96, 3}

```

Figure 12: Wolfram code for recurrence for the variance

4 Basic Kalman Folding

We need matrix operations, now, and we'll use the Gnu Scientific Library, GSL⁹, which includes support for CBLAS¹⁰ and LAPACKE.¹¹

4.1 Get GSL to Work

Get gsl, build it (`./configure,make`), check it (`make check`), install it (`make install`). The following should work if you use all the default settings; you may need to install *gfortran* separately.

You must link these against `libgsl`. Note the flags on the `begin-src` line if you are visiting the `org-mode` file in `emacs`.

```

double x = 5.0;
double y = gsl_sf_bessel_J0 (x);
printf ("J0(%g) = %.18e\n", x, y);
// ~~> produces

J0(5) = -1.775967713143382642e-01

gsl_block * b = gsl_block_alloc (100);

printf ("length of block = %zu\n", b->size);
printf ("block data address = %p\n", b->data);

gsl_block_free (b);
// ~~> produces

length of block = 100
block data address = 0x7fe6d1403580

int i;
gsl_vector * v = gsl_vector_alloc (3);

for (i = 0; i < 3; i++)

```

⁹<http://www.gnu.org/software/gsl/>

¹⁰<http://www.netlib.org/blas/>

¹¹<http://www.netlib.org/lapack/lapacke.html>

```

{    gsl_vector_set (v, i, 1.23 + i);    }

for (i = 0; i < 3; i++)
{    printf ("v_%d = %g\n", i, gsl_vector_get (v, i));    }

gsl_vector_free (v);
// ~~> produces

v0 = 1.23
v1 = 2.23
v2 = 3.23

int i, j;
gsl_matrix * m = gsl_matrix_alloc (10, 3);

for (i = 0; i < 10; i++)
    for (j = 0; j < 3; j++)
        gsl_matrix_set (m, i, j, 0.23 + 100*i + j);

for (i = 0; i < 10; i++)
    for (j = 0; j < 3; j++)
        printf ("m(%d,%d) = %g\n", i, j,
            gsl_matrix_get (m, i, j));

gsl_matrix_free (m);
// ~~> produces

```

```

m(0  0) = 0.23
m(0  1) = 1.23
m(0  2) = 2.23
m(1  0) = 100.23
m(1  1) = 101.23
m(1  2) = 102.23
m(2  0) = 200.23
m(2  1) = 201.23
m(2  2) = 202.23
m(3  0) = 300.23
m(3  1) = 301.23
m(3  2) = 302.23
m(4  0) = 400.23
m(4  1) = 401.23
m(4  2) = 402.23
m(5  0) = 500.23
m(5  1) = 501.23
m(5  2) = 502.23
m(6  0) = 600.23
m(6  1) = 601.23
m(6  2) = 602.23
m(7  0) = 700.23
m(7  1) = 701.23
m(7  2) = 702.23
m(8  0) = 800.23
m(8  1) = 801.23
m(8  2) = 802.23
m(9  0) = 900.23
m(9  1) = 901.23
m(9  2) = 902.23

```

4.2 BLAS Through GSL

You must link this against `libgsl` and `libblas`. Note the flags on the `begin-src` line if you are visiting the org-mode file in emacs.

```

double a[] = { 0.11, 0.12, 0.13,
               0.21, 0.22, 0.23 };

double b[] = { 1011, 1012,
               1021, 1022,
               1031, 1032 };

double c[] = { 0.00, 0.00,
               0.00, 0.00 };

gsl_matrix_view A = gsl_matrix_view_array(a, 2, 3);

```

```

gsl_matrix_view B = gsl_matrix_view_array(b, 3, 2);
gsl_matrix_view C = gsl_matrix_view_array(c, 2, 2);

/* Compute C = A B */

gsl_blas_dgemm (CblasNoTrans, CblasNoTrans,
                1.0, &A.matrix, &B.matrix,
                0.0, &C.matrix);

printf ("[ %g, %g\n", c[0], c[1]);
printf (" %g, %g ]\n", c[2], c[3]);
// ~~> produces

                        [ 367.76   368.12
                        674.06   674.72 ]

```

4.3 Step 2: Getting LAPACK to work

Make sure you have `gfortran` installed:

```
gfortran --version
```

Get LAPACK.¹² This builds BLAS as a side effect.

```

pushd ~/Documents/lapack-3.6.0
cmake .
make
make test

pushd ~/Documents/lapack-3.6.0
make install

```

4.3.1 Make LAPACKE

This is the C interface to LAPACK. The following mercilessly hacks around a couple of problems in the build of `examples`, but it's enough to get the example working.

```

pushd ~/Documents/lapack-3.6.0
cp make.inc.example make.inc
cd LAPACKE
make lapacke

pushd ~/Documents/lapack-3.6.0
find . -name "*.a"

```

¹²<http://www.netlib.org/lapack/>

```
pushd ~/Documents/lapack-3.6.0
cd LAPACKE
cp ./include/lapacke*.h /usr/local/include
cd example
cp ../../liblapacke.a /usr/local/lib
cp ../../lib/*.a ../../
cp ../../libblas.a ../../librefblas.a
make
```

Emacs 24.5.1 (Org mode 8.3.4)